

# pf

Packet Filter pseudo-device

## Name:

/dev/pf

## Description:

Packet filtering takes place in io-pkt. A pseudo-device, /dev/pf, lets user processes control the behavior of the packet filter through an [ioctl\(\)](#) interface. There are commands to enable and disable the filter, load rule sets, add and remove individual rules or state table entries, and retrieve statistics. The most commonly used functions are covered by [pfctl](#).



Although the NetBSD documentation talks about *ioctl()*, you should use *ioctl\_socket()* instead in your packet-filtering code. With the microkernel message-passing architecture, *ioctl()* calls that have pointers embedded in them need to be handled specially. The *ioctl\_socket()* function will default to *ioctl()* for functionality that doesn't require special handling.

Manipulations such as loading a rule set that involve more than a single *ioctl\_socket()* call require *aticket*, which prevents the occurrence of multiple concurrent manipulations.

Fields of *ioctl\_socket()* parameter structures that refer to packet data (such as addresses and ports) are generally expected in network-byte order.

Rules and address tables are contained in *anchors*. When servicing an *ioctl\_socket()* request, if the anchor field of the argument structure is empty, io-pkt uses the default anchor (i.e. the main rule set) in operations. Anchors are specified by name and may be nested, with components separated by slashes, similar to the way that filesystem hierarchies are laid out. The final component of the anchor path is the anchor under which operations will be performed.

## [ioctl\\_socket\(\) interface](#)

The pf pseudo-device supports the following *ioctl\_socket()* commands, available through <net/pfvar.h>:

### DIOCPSTART

Start the packet filter.

### DIOCPSTOP

Stop the packet filter.

### DIOCPSTARTALTQ

Start the ALTQ bandwidth control system (see [altq](#) in the NetBSD documentation).

### DIOCPSTOPALTQ

Stop the ALTQ bandwidth control system.

### DIOCPBEGINADDRS struct pfloc\_pooladdr \*pp

Clear the buffer address pool and get a ticket for subsequent DIOCPADDADDR, DIOCPADDRULE, and DIOCPCHANGERULE calls. The pfloc\_pooladdr structure is defined as follows:

```
struct pfloc_pooladdr {  
  
    u_int32_t          action;  
  
    u_int32_t          ticket;  
  
    u_int32_t          nr;  
  
    u_int32_t          r_num;  
  
    u_int8_t           r_action;  
};
```

```

    u_int8_t            r_last;

    u_int8_t            af;

    char                anchor[MAXPATHLEN];

    struct pf_pooladdr   addr;

};

```

#### **DIOCADDADDR struct pfloc\_pooladdr \*pp**

Add the pool address, *addr* to the buffer address pool to be used in the following DIOCADDRULE or DIOCCHANGERULE call. All other members of the structure are ignored.

#### **DIOCADDRULE struct pfloc\_rule \*pr**

Add the given rule at the end of the inactive rule set. The *pfloc\_rule* structure is defined as follows:

```

struct pfloc_rule {

    u_int32_t           action;

    u_int32_t           ticket;

    u_int32_t           pool_ticket;

    u_int32_t           nr;

    char                anchor[MAXPATHLEN];

    char                anchor_call[MAXPATHLEN];

    struct pf_rule       rule;

};

```

This call requires a *ticket* obtained through a preceding DIOCBEGIN call and a *pool\_ticket* obtained through a DIOCBEGINADDRS call. You must also call DIOCADDADDR if any pool addresses are required.

The optional anchor name indicates the anchor in which to append the rule. The *nr* and *action* members are ignored.

#### **DIOCADDALTQ struct pfloc\_altq \*pa**

Add an ALTQ discipline or queue. The *pfloc\_altq* structure is defined as follows:

```

struct pfloc_altq {

    u_int32_t           action;

    u_int32_t           ticket;

};

```

```

        u_int32_t        nr;

        struct pf_altq   altq;

};

```

#### **DIOCGETRULES struct pfloc\_rule \*pr**

Get a ticket for subsequent DIOCGETRULE calls, and the number *nr* of rules in the active rule set.

#### **DIOCGETRULE struct pfloc\_rule \*pr**

Get a rule by its number *nr*, using the ticket obtained through a preceding DIOCGETRULES call.

#### **DIOCGETADDRS struct pfloc\_pooladdr \*pp**

Get a ticket for subsequent DIOCGETADDR calls and the number *nr* of pool addresses in the rule specified with *r\_action*, *r\_num*, and *anchor*.

#### **DIOCGETADDR struct pfloc\_pooladdr \*pp**

Get the pool address *addr* by its number *nr* from the rule specified with *r\_action*, *r\_num*, and *anchor*, using the ticket obtained through a preceding DIOCGETADDRS call.

#### **DIOCGETALTQS struct pfloc\_altq \*pa**

Get a ticket for subsequent DIOCGETALTQ calls and the number *nr* of queues in the active list.

#### **DIOCGETALTQ struct pfloc\_altq \*pa**

Get the queueing discipline *altq* by its number *nr*, using the ticket obtained through a preceding DIOCGETALTQS call.

#### **DIOCGETQSTATS struct pfloc\_qstats \*pq**

Get the statistics for a queue. The *pfloc\_qstats* structure is defined as follows:

```

struct pfloc_qstats {

        u_int32_t        ticket;

        u_int32_t        nr;

        void             *buf;

        int nbytes;

        u_int8_t         scheduler;

};

```

This call fills in a pointer to the buffer of statistics *buf*, of length *nbytes*, for the queue specified by *nr*.

#### **DIOCGETRULESETS struct pfloc\_ruleset \*pr**

Get the number *nr* of rule sets (i.e., anchors) directly attached to the anchor named by *path* for use in subsequent DIOCGETRULESET calls. The *pfloc\_ruleset* structure is defined as follows:

```

struct pfloc_ruleset {

```

```

        u_int32_t      nr;

        char           path[MAXPATHLEN];

        char           name[PF_ANCHOR_NAME_SIZE];

};

```

Nested anchors, since they aren't directly attached to the given anchor, aren't included. This *ioctl\_socket()* command returns EINVAL if the given anchor doesn't exist.

#### **DIOCGRULESET struct pfloc\_ruleset \*pr**

Get a rule set (i.e., an anchor) name by its number *nr* from the given anchor path, the maximum number of which can be obtained from a preceding DIOCGRULESETS call. This *ioctl\_socket()* command returns EINVAL if the given anchor doesn't exist, or EBUSY if another process is concurrently updating a rule set.

#### **DIOCADDSTATE struct pfloc\_state \*ps**

Add a state entry. The pfloc\_state structure is defined as follows:

```

struct pfloc_state {

        u_int32_t      nr;

        struct pf_state state;

};

```

#### **DIOCGETSTATE struct pfloc\_state \*ps**

Extract the entry with the specified number *nr* from the state table.

#### **DIOCKILLSTATES struct pfloc\_state\_kill \*psk**

Remove matching entries from the state table. This *ioctl\_socket()* command returns the number of killed states in *psk\_af*. The pfloc\_state\_kill structure is defined as follows:

```

struct pfloc_state_kill {

        sa_family_t      psk_af;

        int               psk_proto;

        struct pf_rule_addr psk_src;

        struct pf_rule_addr psk_dst;

        char              psk_ifname[IFNAMSIZ];

};

```

### **DIOCCLRSTATES struct pfloc\_state\_kill \*psk**

Clear all states. This command works like DIOCKILLSTATES, but ignores the *psk\_af*, *psk\_proto*, *psk\_src*, and *psk\_dst* fields of the *pfloc\_state\_kill* structure.

### **DIOCSETSTATUSIF struct pfloc\_if \*pi**

Specify the interface for which to gather statistics. The *pfloc\_if* structure is defined as follows:

```
struct pfloc_if {  
  
    char            ifname[IFNAMSIZ];  
  
};
```

### **DIOCGETSTATUS struct pf\_status \*s**

Get the internal packet filter statistics. The *pf\_status* structure is defined as follows:

```
struct pf_status {  
  
    u_int64_t        counters[PFRES_MAX];  
  
    u_int64_t        lcounters[LCNT_MAX];  
  
    u_int64_t        fcounters[FCNT_MAX];  
  
    u_int64_t        scounters[SCNT_MAX];  
  
    u_int64_t        pcounters[2][2][3];  
  
    u_int64_t        bcounters[2][2];  
  
    u_int64_t        stateid;  
  
    u_int32_t        running;  
  
    u_int32_t        states;  
  
    u_int32_t        src_nodes;  
  
    u_int32_t        since;  
  
    u_int32_t        debug;  
  
    u_int32_t        hostid;  
  
    char            ifname[IFNAMSIZ];  
};
```

```
};
```

### **DIOCCLRSTATUS**

Clear the internal packet filter statistics.

### **DIOCSTATLOOK struct pfloc\_natlook \*pnl**

Look up a state table entry by source and destination addresses and ports. The pfloc\_natlook structure is defined as follows:

```
struct pfloc_natlook {  
  
    struct pf_addr    saddr;  
  
    struct pf_addr    daddr;  
  
    struct pf_addr    rsaddr;  
  
    struct pf_addr    rdaddr;  
  
    u_int16_t         sport;  
  
    u_int16_t         dport;  
  
    u_int16_t         rsport;  
  
    u_int16_t         rdport;  
  
    sa_family_t       af;  
  
    u_int8_t          proto;  
  
    u_int8_t          direction;  
  
};
```

### **DIOCSETDEBUG u\_int32\_t \*level**

Set the debug level to one of PF\_DEBUG\_NONE, PF\_DEBUG\_URGENT, PF\_DEBUG\_MISC, or PF\_DEBUG\_NOISY.

### **DIOCGETSTATES struct pfloc\_states \*ps**

Get state table entries. The pfloc\_states structure is defined as follows:

```
struct pfloc_states {  
  
    int        ps_len;  
  
    union {
```

```

        caddr_t        psu_buf;

        struct pf_state *psu_states;

    } ps_u;

#define ps_buf        ps_u.psu_buf

#define ps_states      ps_u.psu_states

};

```

If *ps\_len* is zero, all states are gathered into *pf\_states*, and *ps\_len* is set to the size they take in memory (i.e. `sizeof(struct pf_state) * nr`). If *ps\_len* is nonzero, as many states that can fit into *ps\_len* as possible are gathered, and *ps\_len* is updated to the size those rules take in memory.

#### **DIOCCHANGERULE struct pfloc\_rule \*pcr**

Add or remove the rule in the rule set specified by *rule.action*. The type of operation to be performed is indicated by *action*, which can be any of the following:

- PF\_CHANGE\_NONE
- PF\_CHANGE\_ADD\_HEAD
- PF\_CHANGE\_ADD\_TAIL
- PF\_CHANGE\_ADD\_BEFORE
- PF\_CHANGE\_ADD\_AFTER
- PF\_CHANGE\_REMOVE
- PF\_CHANGE\_GET\_TICKET

You must set *ticket* to the value obtained with PF\_CHANGE\_GET\_TICKET for all actions except PF\_CHANGE\_GET\_TICKET. You must set *pool\_ticket* to the value obtained with the DIOCBEGINADDRS call for all actions except PF\_CHANGE\_REMOVE and PF\_CHANGE\_GET\_TICKET. The *anchor* indicates which anchor the operation applies to. The *nr* member indicates the rule number against which to apply PF\_CHANGE\_ADD\_BEFORE, PF\_CHANGE\_ADD\_AFTER, or PF\_CHANGE\_REMOVE actions.

#### **DIOCCHANGEADDR struct pfloc\_pooladdr \*pca**

Add or remove the pool address *addr* from the rule specified by *r\_action*, *r\_num*, and *anchor*.

#### **DIOCSETTIMEOUT struct pfloc\_tm \*pt**

Set the state timeout of *timeout* to *seconds*. The *pfloc\_tm* structure is defined as follows:

```

struct pfloc_tm {

    int timeout;

    int seconds;

};

```

The old value is placed into *seconds*. For the possible values of *timeout*, see the PFTM\_\* values in `<net/pfvar.h>`.

#### **DIOCGETTIMEOUT struct pfloc\_tm \*pt**

Get the state timeout of *timeout*. The value is placed into the *seconds* field.

#### **DIOCCLRRELECTRS**

Clear per-rule statistics.

#### **DIOCSETLIMIT struct pfloc\_limit \*pl**

Set the hard limits on the memory pools used by the packet filter. The *pfloc\_limit* structure is defined as follows:

```
struct pfloc_limit {

    int          index;

    unsigned     limit;

};

enum { PF_LIMIT_STATES, PF_LIMIT_SRC_NODES,

       PF_LIMIT_FRAGS };
```

#### **DIOCGETLIMIT struct pfloc\_limit \*pl**

Get the hard limit for the memory pool indicated by *index*.

#### **DIOCRCLRTABLES struct pfloc\_table \*io**

Clear all tables. All the *ioctl\_socket()* commands that manipulate radix tables use the same structure described below. On exit from the DIOCRCLRTABLES command, *pfrio\_ndel* contains the number of tables deleted.

The *pfloc\_table* structure is defined as follows:

```
struct pfloc_table {

    struct pfr_table      pfrio_table;

    void                  *pfrio_buffer;

    int                   pfrio_esize;

    int                   pfrio_size;

    int                   pfrio_size2;

    int                   pfrio_nadd;

    int                   pfrio_ndel;

    int                   pfrio_nchange;

    int                   pfrio_flags;

    u_int32_t             pfrio_ticket;
```



```
};

#define pfrio_exists    pfrio_nadd

#define pfrio_nzero     pfrio_nadd

#define pfrio_nmatch    pfrio_nadd

#define pfrio_naddr     pfrio_size2

#define pfrio_setflag   pfrio_size2

#define pfrio_clrflag   pfrio_nadd
```

#### **DIOCRADDTABLES struct pfloc\_table \*io**

Create one or more tables. On entry, *pfrio\_buffer[pfrio\_size]* contains a table of pfr\_table structures. On exit, *pfrio\_nadd* contains the number of tables effectively created. The pfr\_table structure is defined as follows:

```
struct pfr_table {

    char            pfrt_anchor[MAXPATHLEN];

    char            pfrt_name[PF_TABLE_NAME_SIZE];

    u_int32_t       pfrt_flags;

    u_int8_t        pfrt_fback;

};
```

#### **DIOCRDELTABLES struct pfloc\_table \*io**

Delete one or more tables. On entry, *pfrio\_buffer[pfrio\_size]* contains a table of pfr\_table structures. On exit, *pfrio\_nadd* contains the number of tables effectively deleted.

#### **DIOCRGETTABLES struct pfloc\_table \*io**

Get a list of all tables. On entry, *pfrio\_buffer[pfrio\_size]* contains a valid writeable buffer for pfr\_table structures. On exit, *pfrio\_size* contains the number of tables written into the buffer. If the buffer is too small, io-pkt doesn't store anything, but just returns the required buffer size, without error.

#### **DIOCRGETTSTATS struct pfloc\_table \*io**

This call is like DIOCRGETTABLES, but is used to get an array of pfr\_tstats structures. The pfr\_tstats structure is defined as follows:

```
struct pfr_tstats {
```

```

    struct pfr_table pfrts_t;

    u_int64_t        pfrts_packets

                        [PFR_DIR_MAX][PFR_OP_TABLE_MAX];

    u_int64_t        pfrts_bytes

                        [PFR_DIR_MAX][PFR_OP_TABLE_MAX];

    u_int64_t        pfrts_match;

    u_int64_t        pfrts_nomatch;

    long             pfrts_tzero;

    int              pfrts_cnt;

    int              pfrts_refcnt[PFR_REFCNT_MAX];

};

#define pfrts_name        pfrts_t.pfrt_name

#define pfrts_flags        pfrts_t.pfrt_flags

```

#### **DIOCRCLRTSTATS struct pfio\_table \*io**

Clear the statistics of one or more tables. On entry, *pfrio\_buffer[pfrio\_size]* contains a table of *pfr\_table* structures. On exit, *pfrio\_nzero* contains the number of tables effectively cleared.

#### **DIOCRCLRADDRS struct pfio\_table \*io**

Clear all addresses in a table. On entry, *pfrio\_table* contains the table to clear. On exit, *pfrio\_ndel* contains the number of addresses removed.

#### **DIOCRADDADDRS struct pfio\_table \*io**

Add one or more addresses to a table. On entry, *pfrio\_table* contains the table ID, and *pfrio\_buffer[pfrio\_size]* contains the list of *pfr\_addr* structures to add. On exit, *pfrio\_nadd* contains the number of addresses effectively added.

The *pfr\_addr* structure is defined as follows:

```

struct pfr_addr {

    union {

        struct in_addr    _pfra_ip4addr;

        struct in6_addr   _pfra_ip6addr;

    }    pfra_u;

```



```

        long                pfras_tzero;

};

```

#### **DIOCRCLRASTATS struct pfloc\_table \*io**

Clear the statistics of one or more addresses. On entry, *pfrio\_table* contains the table ID, and *pfrio\_buffer[pfrio\_size]* contains a table of pfr\_addr structures to clear. On exit, *pfrio\_nzero* contains the number of addresses effectively cleared.

#### **DIOCRTSTADDRS struct pfloc\_table \*io**

Test if the given addresses match a table. On entry, *pfrio\_table* contains the table ID, and *pfrio\_buffer[pfrio\_size]* contains a table of pfr\_addr structures to test. On exit, io-pkt updates the *pfr\_addr* table by setting the *pfr\_fback* member appropriately.

#### **DIOCRSETTFLAGS struct pfloc\_table \*io**

Change the PFR\_TFLAG\_CONST or PFR\_TFLAG\_PERSIST flags of a table. On entry, *pfrio\_buffer[pfrio\_size]* contains a table of pfr\_table structures, and *pfrio\_setflag* contains the flags to add, while *pfrio\_clrflag* contains the flags to remove. On exit, *pfrio\_nchange* and *pfrio\_ndel* contain the number of tables altered or deleted by io-pkt.



You can delete tables if you remove the PFR\_TFLAG\_PERSIST flag of an unreferenced table.

#### **DIOCRINADEFINE struct pfloc\_table \*io**

Define a table in the inactive set. On entry, *pfrio\_table* contains the table ID, and *pfrio\_buffer[pfrio\_size]* contains the list of pfr\_addr structures to put in the table. A valid ticket must also be supplied to *pfrio\_ticket*. On exit, *pfrio\_nadd* contains 0 if the table was already defined in the inactive list, or 1 if a new table has been created. The *pfrio\_naddr* member contains the number of addresses effectively put in the table.

#### **DIOCXBEGIN struct pfloc\_trans \*io**

Clear all the inactive rule sets specified in the *pfloc\_trans\_e* array. The pfloc\_trans structure is defined as follows:

```

struct pfloc_trans {

    int    size; /* number of elements */

    int    esize; /* size of each element in bytes */

    struct pfloc_trans_e {

        int            rs_num;

        char            anchor[MAXPATHLEN];

        u_int32_t       ticket;

    } *array;

};

```

For each rule set, a ticket is returned for subsequent “add rule” *ioctl\_socket()* commands, as well as for the DIOCXCOMMIT and DIOXROLLBACK calls.

Rule set types, identified by *rs\_num*, include the following:

- PF\_RULESET\_SCRUB — scrub (packet normalization) rules.
- PF\_RULESET\_FILTER — filter rules.
- PF\_RULESET\_NAT — NAT (Network Address Translation) rules.
- PF\_RULESET\_BINAT — bidirectional NAT rules.
- PF\_RULESET\_RDR — redirect rules.
- PF\_RULESET\_ALTQ — ALTQ disciplines.
- PF\_RULESET\_TABLE — address tables.

#### **DIOXCOMMIT struct pfloc\_trans \*io**

Atomically switch a vector of inactive rule sets to the active rule sets. This call is implemented as a standard two-phase commit, which either fails for all rule sets, or completely succeeds. All tickets need to be valid.

This *ioctl\_socket()* command returns EBUSY if another process is concurrently updating some of the same rule sets.

#### **DIOXROLLBACK struct pfloc\_trans \*io**

Clean up io-pkt by undoing all changes that have taken place on the inactive rule sets since the lastDIOXBEGIN. DIOXROLLBACK silently ignores rule sets for which the ticket is invalid.

#### **DIOXHOSTID u\_int32\_t \*hostid**

Set the host ID, which is used by pfsync to identify which host created state table entries.

#### **DIOXSPFLUSH**

Flush the passive OS fingerprint table.

#### **DIOXSPADD struct pf\_osfp\_ioctl \*io**

Add a passive OS fingerprint to the table. The pf\_osfp\_ioctl structure is defined as follows:

```
struct pf_osfp_ioctl {

    struct pf_osfp_entry {

        SLIST_ENTRY(pf_osfp_entry) fp_entry;

        pf_osfp_t    fp_os;

        char          fp_class_nm[PF_OSFP_LEN];

        char          fp_version_nm[PF_OSFP_LEN];

        char          fp_subtype_nm[PF_OSFP_LEN];

    } fp_os;

    pf_tcpopts_t      fp_tcpopts;

    u_int16_t         fp_wsize;

    u_int16_t         fp_psize;
```

```

    u_int16_t      fp_mss;

    u_int16_t      fp_flags;

    u_int8_t       fp_optcnt;

    u_int8_t       fp_wscale;

    u_int8_t       fp_ttl;

    int            fp_getnum;

};

```

Set *fp\_os.fp\_os* to the packed fingerprint, *fp\_os.fp\_class\_nm* to the name of the class (Linux, Windows, etc.), *fp\_os.fp\_version\_nm* to the name of the version (NT, 95, 98), and *fp\_os.fp\_subtype\_nm* to the name of the subtype or patch level. The members *fp\_mss*, *fp\_wsize*, *fp\_psize*, *fp\_ttl*, *fp\_optcnt*, and *fp\_wscale* are set to the TCP MSS, the TCP window size, the IP length, the IP TTL, the number of TCP options, and the TCP window scaling constant of the TCP SYN packet, respectively.

The *fp\_flags* member is filled according to the PF\_OSFP\_\* definition in <net/pfvar.h>. The *fp\_tcpopts* member contains packed TCP options. Each option uses PF\_OSFP\_TCPOPT\_BITS bits in the packed value. Options include any of PF\_OSFP\_TCPOPT\_NOP, PF\_OSFP\_TCPOPT\_SACK, PF\_OSFP\_TCPOPT\_WSCALE, PF\_OSFP\_TCPOPT\_MSS, or PF\_OSFP\_TCPOPT\_TS.

This *ioctl\_socket()* command doesn't use the *fp\_getnum* member.



You must zero the structure's slack space for correct operation; *memset()* the whole structure to zero before filling and sending it to *io-pkt*.

#### **DIOSCOSFPGET struct pf\_osfp\_ioctl \*io**

Get the passive OS fingerprint number *fp\_getnum* from *io-pkt*'s fingerprint list. The rest of the structure members comes back filled. Get the whole list by repeatedly incrementing the *fp\_getnum* number until *ioctl\_socket()* gives an error of EBUSY.

#### **DIOCGETSRCNODES struct pfloc\_src\_nodes \*psn**

Get the list of source nodes kept by sticky addresses and source tracking. The *pfloc\_src\_nodes* structure is defined as follows:

```

struct pfloc_src_nodes {

    int      psn_len;

    union {

        caddr_t      psu_buf;

        struct pf_src_node      *psu_src_nodes;
    };
};

```

```

    } psn_u;

#define psn_buf          psn_u.psu_buf

#define psn_src_nodes    psn_u.psu_src_nodes

};

```

You must call *ioctl\_socket()* once with *psn\_len* set to 0. If *ioctl\_socket()* returns without error, *psn\_len* is set to the size of the buffer required to hold all the *pf\_src\_node* structures held in the table. You should then allocate a buffer of this size and place a pointer to this buffer in *psn\_buf*. You must then call *ioctl\_socket()* again to fill this buffer with the actual source node data. After that call, *psn\_len* is set to the length of the buffer actually used.

### DIOCLRSRCNODES

Clear the tree of source-tracking nodes.

### DIOCIGETIFACES struct pfloc\_iface \*io

Get a list of interfaces and interface drivers known to pf. All the *ioctl\_socket()* commands that manipulate interfaces use the structure described below:

```

struct pfloc_iface {

    char                pflio_name[IFNAMSIZ];

    void                *pflio_buffer;

    int                 pflio_esize;

    int                 pflio_size;

    int                 pflio_nzero;

    int                 pflio_flags;

};

#define PFI_FLAG_GROUP    0x0001  /* gets groups of interfaces */

#define PFI_FLAG_INSTANCE 0x0002  /* gets single interfaces */

#define PFI_FLAG_ALLMASK  0x0003

```

If it isn't empty, you can use *pflio\_name* to restrict the search to a specific interface or driver.

The *pflio\_buffer[pflio\_size]* member is the user-supplied buffer for returning the data. On entry, *pflio\_size* represents

the number of `pfi_if` entries that can fit into the buffer. The io-pkt manager replaces this value with the real number of entries it wants to return.

You should set `pfio_esize` to `sizeof(struct pfi_if)`. You should set `pfio_flags` to `PFI_FLAG_GROUP`, `PFI_FLAG_INSTANCE`, or both, to tell io-pkt to return a group of interfaces (drivers, such as `fxp`), real interface instances (e.g. `fxp1`), or both. The data is returned in the `pfi_if` structure described below:

```
struct pfi_if {

    char          pfif_name[IFNAMSIZ];

    u_int64_t     pfif_packets[2][2][2];

    u_int64_t     pfif_bytes[2][2][2];

    u_int64_t     pfif_addcnt;

    u_int64_t     pfif_delcnt;

    long          pfif_tzero;

    int           pfif_states;

    int           pfif_rules;

    int           pfif_flags;

};

#define PFI_IFLAG_GROUP      0x0001  /* group of interfaces */

#define PFI_IFLAG_INSTANCE  0x0002  /* single instance */

#define PFI_IFLAG_CLONABLE  0x0010  /* clonable group */

#define PFI_IFLAG_DYNAMIC   0x0020  /* dynamic group */

#define PFI_IFLAG_ATTACHED  0x0040  /* interface attached */
```

#### **DIOICLRISTATS struct pfloc\_iface \*io**

Clear the statistics counters of one or more interfaces. You can use `pfio_name` and `pfio_flags` to select which interfaces need to be cleared. The filtering process is the same as for `DIOCGETIFACES`. The `pfio_nzero` member is set by io-pkt to the number of interfaces and drivers that have been cleared.

#### **DIOCSETIFFLAG struct pfloc\_iface \*io**

Set the user-setable flags of the pf internal interface description:

- `PFI_IFLAG_SKIP` — skip the interface.



The filtering process is the same as for DIOCGETIFACES.

### **DIOCCLRIFFLAG struct pfloc\_iface \*io**

Similar to DIOCSETIFFLAG, but clear the flags.

### **Examples:**

The following example demonstrates how to use the DIOCNATLOOK command to find the internal host/port of a NATed connection:

```
#include <sys/types.h>

#include <sys/socket.h>

#include <sys/ioctl.h>

#include <sys/fcntl.h>

#include <net/if.h>

#include <netinet/in.h>

#include <net/pfvar.h>

#include <err.h>

#include <stdio.h>

#include <stdlib.h>

u_int32_t

read_address(const char *s)

{

    int a, b, c, d;

    sscanf(s, "%i.%i.%i.%i", &a, &b, &c, &d);

    return htonl(a << 24 | b << 16 | c << 8 | d);

}

void

print_address(u_int32_t a)
```

```

{

    a = ntohl(a);

    printf("%d.%d.%d.%d", a >> 24 & 255, a >> 16 & 255,

        a >> 8 & 255, a & 255);

}

int

main(int argc, char *argv[])

{

    struct pfio natlook nl;

    int dev;

    if (argc != 5) {

        printf("%s <gw addr> <gw port> <ext addr> <ext port>\n",

            argv[0]);

        return 1;

    }

    dev = open("/dev/pf", O_RDWR);

    if (dev == -1)

        err(1, "open(\"/dev/pf\") failed");

    memset(&nl, 0, sizeof(struct pfio natlook));

    nl.saddr.v4.s_addr = read_address(argv[1]);

    nl.sport = htons(atoi(argv[2]));

```

```

    nl.daddr.v4.s_addr      = read_address(argv[3]);

    nl.dport                = htons(atoi(argv[4]));

    nl.af                   = AF_INET;

    nl.proto                 = IPPROTO_TCP;

    nl.direction            = PF_IN;


    if (ioctl_socket(dev, DIOCNATLOOK, &nl))

        err(1, "DIOCNATLOOK");


    printf("internal host ");

    print_address(nl.rsaddr.v4.s_addr);

    printf(":%u\n", ntohs(nl.rsport));

    return 0;

}

```

### **Caveats:**

The following functionality is missing from pf in this version of NetBSD:

- The pfsync protocol isn't supported.
- The group keyword isn't supported.

### **See also:**

[pfctl](#)

[ioctl\(\)](#) in the Neutrino Library Reference